
Lecture 6

Deadlocks

I233E OPERATING SYSTEMS

RAZVAN BEURAN

Today's Topics

Deadlocks

- What they are

Mechanisms

- Deadlock prevention
- Deadlock avoidance
- Deadlock detection

Deadlock Model

Resource type

- CPU, memory location, file, I/O device, ...

Resource instance

- One resource item of a given type

Steps to use the resource

- Processes request some **resource type**
- Request is fulfilled with **any instance** of the requested type
- Processes **use** that resource instance, then **release** it

Sequence: *Request* → *Use* → *Release*

Necessary Conditions for Deadlock

Mutual exclusion

- Some resource is used exclusively

Hold and wait

- Some process holds a resource and waits for others

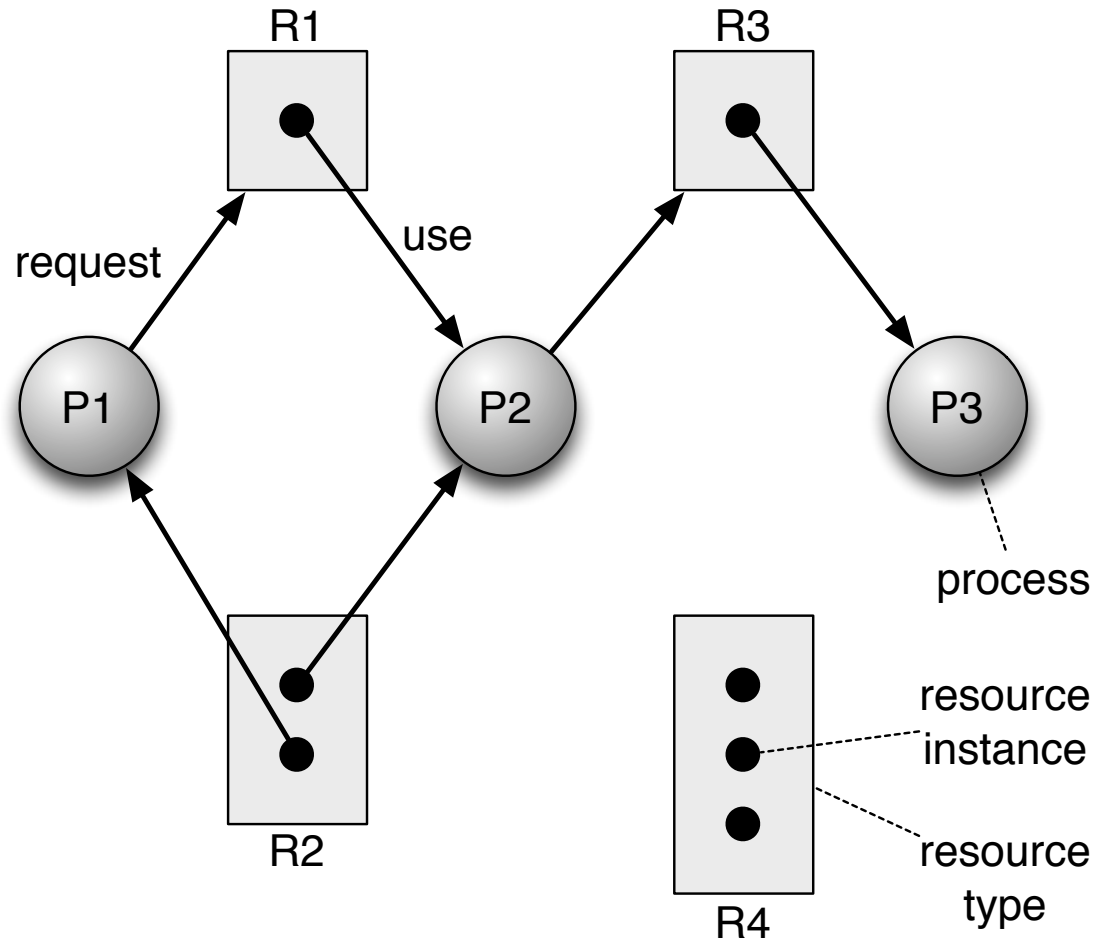
No preemption

- Resource can only be released voluntarily

Circular wait

- Cycle: P_0 waits for P_1 , which waits for ..., which waits for P_0

Resource Allocation Graph



Graph Interpretation

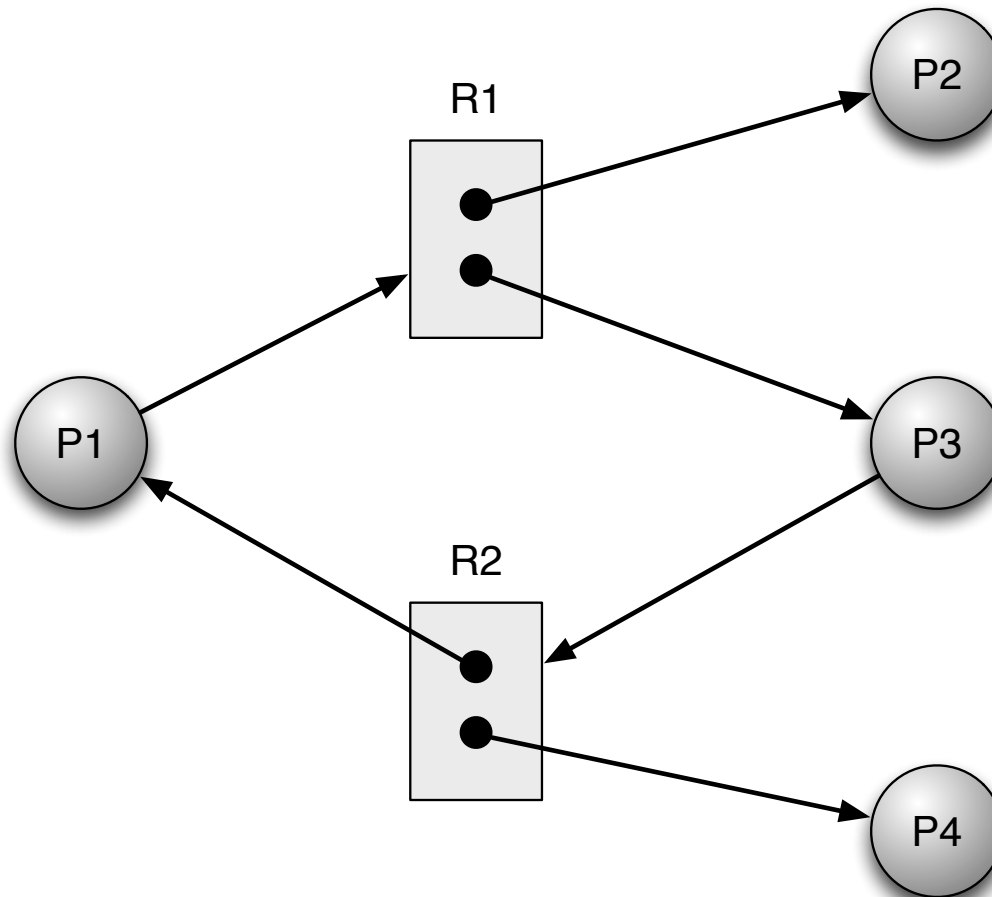
General case

- No cycle → no deadlock can occur
- Cycle means that it is *possible* for a deadlock to occur

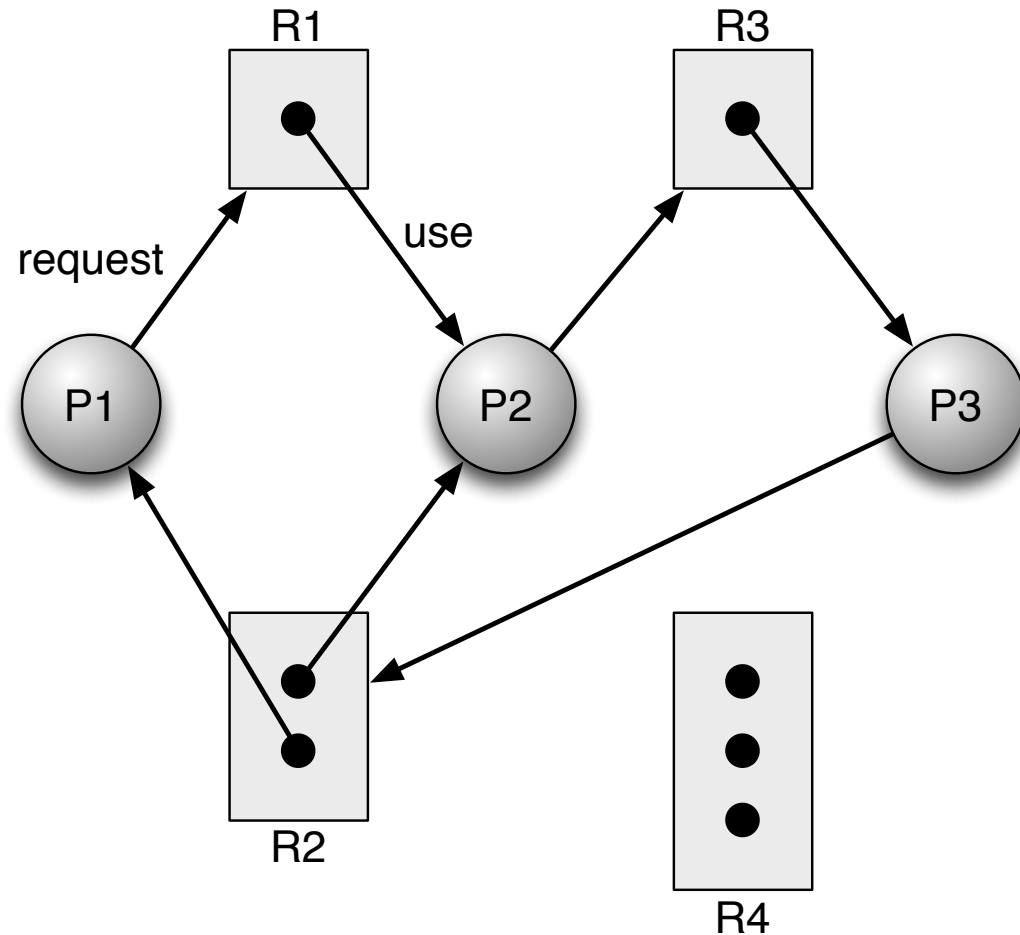
Special case: 1 resource type, 1 resource instance

- Cycle means **deadlock**

Cycle, but No Deadlock



Deadlock Example



Dealing With Deadlocks

Approach 1

- **Prevent** / **avoid** deadlocks
 - Prevent = Remove one condition for deadlock
 - Avoid = Keep out of “unsafe” situations

Approach 2

- When deadlock occurs
 - **Detect** the deadlock
 - **Recover** from deadlock

Approach 3

- Ignore the problem... and pray 😊

Deadlock Prevention

Prevention Mechanisms

Idea: Avoid any of the 4 necessary deadlock conditions

Avoid mutual exclusion →

- Allow resources to be shared

Avoid hold and wait →

- Make all requests first, before execution (*grouped request*)

Avoid no preemption →

- Allow forced release of resources by OS

Avoid circular wait →

- Create total ordering on resources, e.g., request resources in *increasing order* of id

Deadlock Avoidance

Avoidance Mechanisms

Idea

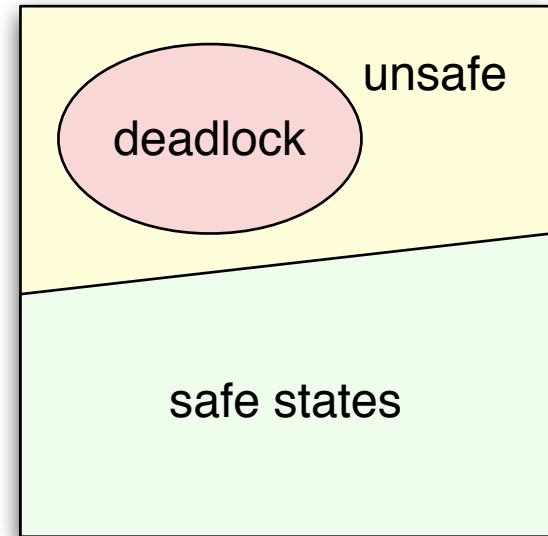
- Get additional information about requests
- Keep track of resource allocation state
- Allow requests only if they are “safe”

Safe state

- There exists some allocation sequence that leads to normal termination

Unsafe state

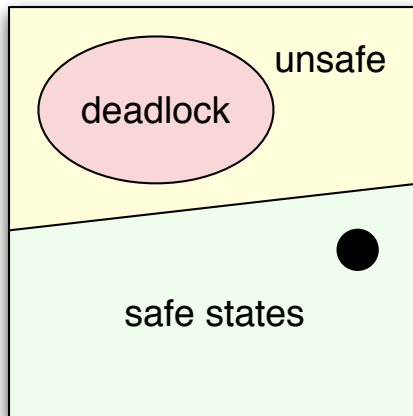
- Every allocation sequence leads to deadlock



Safe vs. Unsafe Allocation (1)

Proc. ID	Max. Needs	Current Holds
P0	10	5
P1	4	2
P2	9	2

Total resource instances = 12



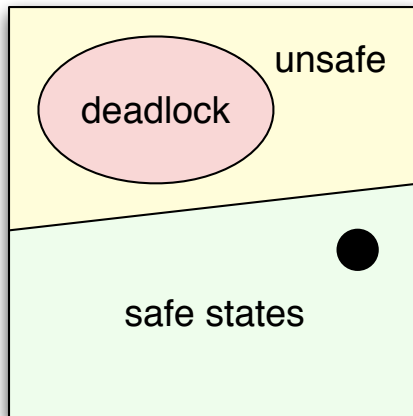
Allocation Sequence #1

Time	P0	P1	P2	Free
t0	5	2	2	3
t1	5	4	2	1
t2	5	-	2	5
t3	10	-	2	0
t4	-	-	2	10
t5	-	-	9	3
t6	-	-	-	12

Safe vs. Unsafe Allocation (2)

Proc. ID	Max. Needs	Current Holds
P0	10	5
P1	4	2
P2	9	2

Total resource instances = 12



Allocation Sequence #2

Time	P0	P1	P2	Free
t0	5	2	2	3
t1	5	2	3	2
t2	5	4	3	0
t3	5	-	3	4
t4	5; req: 5	-	3; req: 6	4

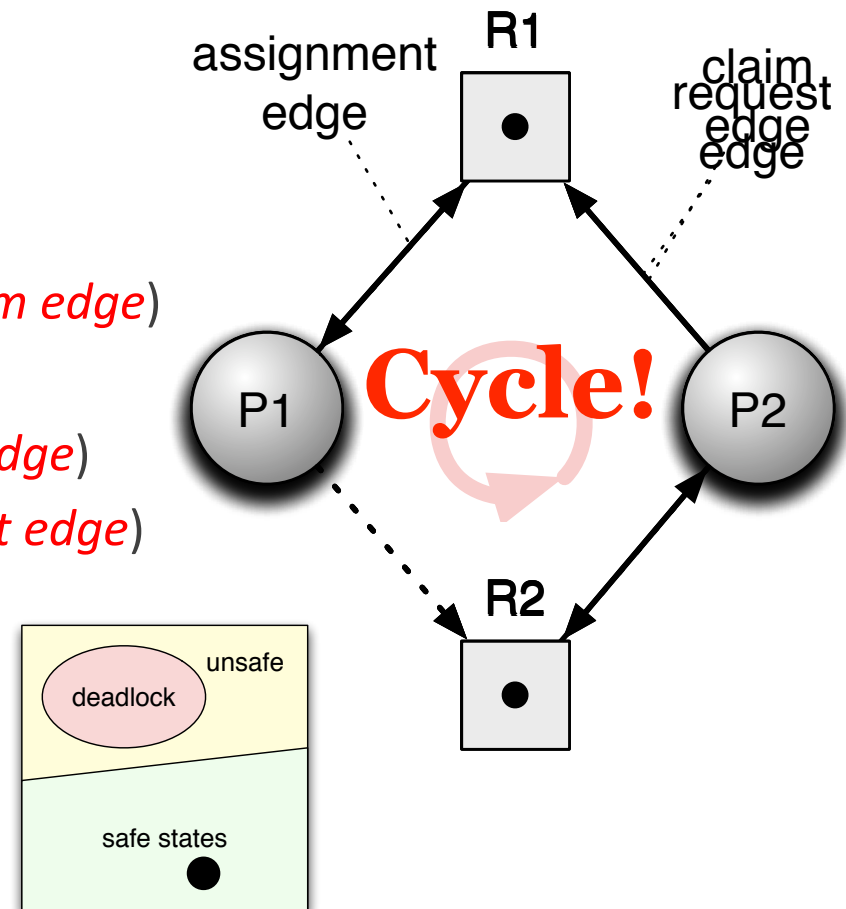


Single-Instance Case

Allocation graph algorithm

- Initially
 - Process “claims” resources (*claim edge*)
- Runtime
 - Resource is requested (*request edge*)
 - Resource is assigned (*assignment edge*)

Unsafe when cycle is created



Banker's Algorithm

System assumptions

- Several resource types
- Several resource instances of each type
- Initially, processes announce max. need for each type

Has 2 parts

- **Safety algorithm**
 - Detect if current system state is safe
- **Resource request algorithm**
 - Authorize (or not) resource requests

Banker's Algorithm (cont.)

System assumptions

- m resource types; n processes

Data structures

- **Available** (vector of size m)
 - $Available[j] = \#$ available instances of resource type j
- **Max** ($n \times m$ matrix)
 - $Max[i,j] = \text{max. demand of } P_i \text{ for resource type } j$
- **Allocation** ($n \times m$ matrix)
 - $Allocation[i,j] = \#$ instances of type j held by P_i
- **Need** ($n \times m$ matrix)
 - $Need[i,j] = Max[i,j] - Allocation[i,j]$

Safety Algorithm

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Safety Algorithm (2)

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	4	3	3	3	2
P_1	2	0	0	1	2	2			
P_2	3	0	2	6	0	0			
P_3	2	1	1	0	1	1			
P_4	0	0	2	4	3	1			

$Need := Max - Allocation$

Safety Algorithm (3)

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Finish		i
P0	x	<input type="checkbox"/>
P1	x	
P2	x	
P3	x	
P4	x	
		Work
		A B C
		3 3 2

Safety Algorithm (4)

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Finish		i
P0	×	1
P1	×	
P2	×	Work
P3	×	A B C
P4	×	3 3 2

Safety Algorithm (5)

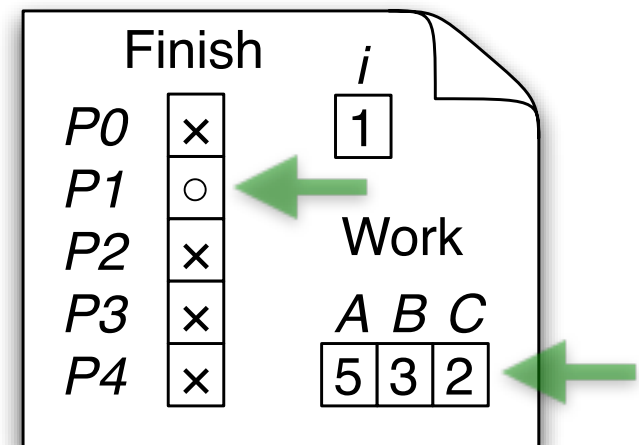
Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			



Safety Algorithm (6)

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

Finish		i
P0	×	3
P1	○	
P2	×	
P3	○	
P4	×	

Work	
A B C	
7 4 3	

Safety Algorithm (7)

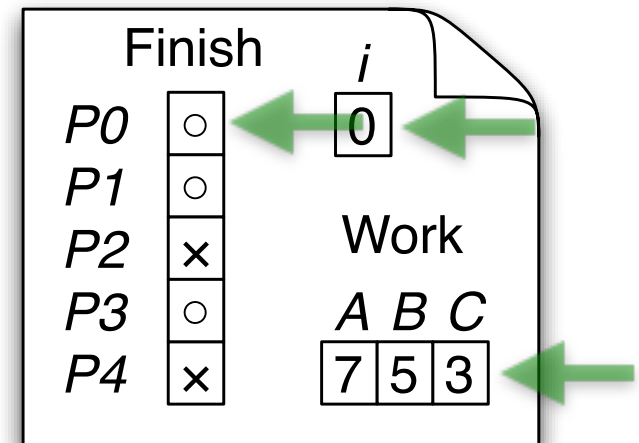
Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
✓ P0	0	1	0	7	4	3	3	3	2
✓ P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
✓ P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			



Safety Algorithm (8)

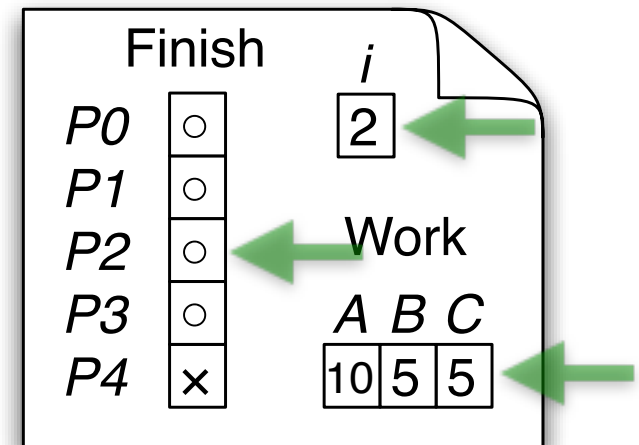
Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
✓ P0	0	1	0	7	4	3	3	3	2
✓ P1	2	0	0	1	2	2			
✓ P2	3	0	2	6	0	0			
✓ P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			



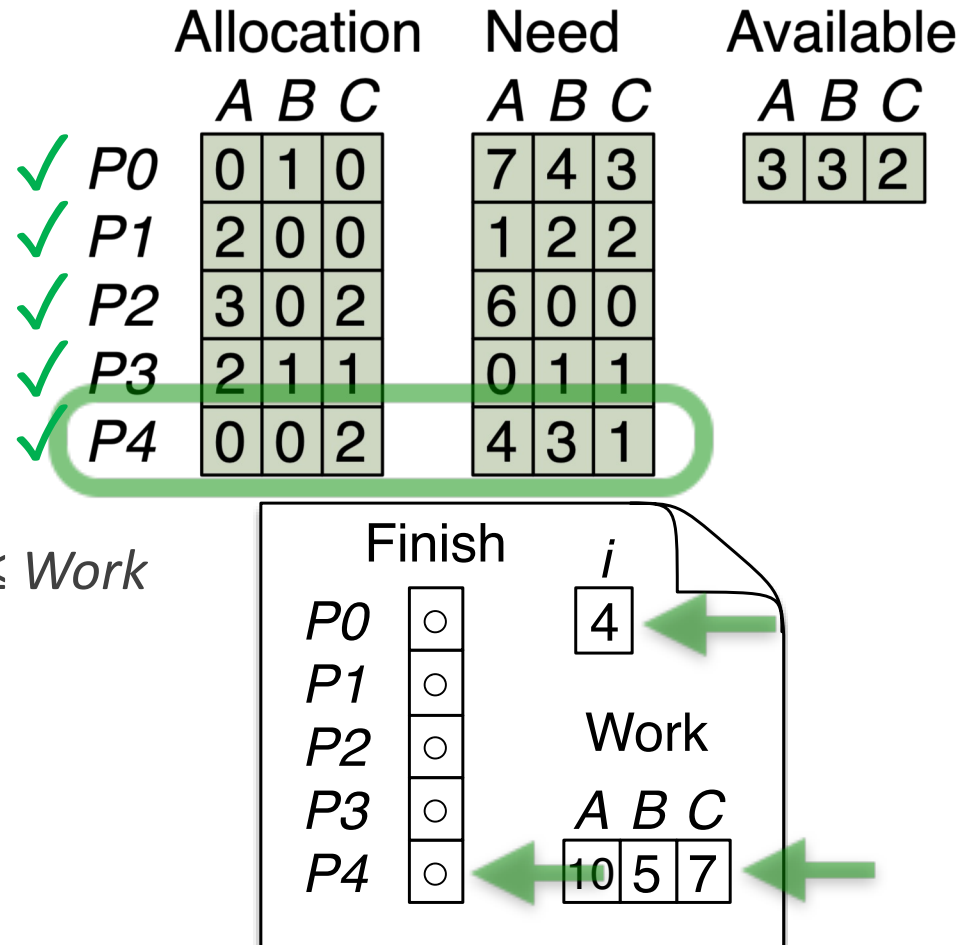
Safety Algorithm (9)

Variables

- Define $Work := Available$
- Define $Finish[0...n-1] := false$

Steps

- While exists i s.t.
 $Finish[i] == false \ \&\& \ Need[i] \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$



Safety Algorithm (10)

Variables

- Define $Work := Available$
- Define $Finish[0..n-1] := false$

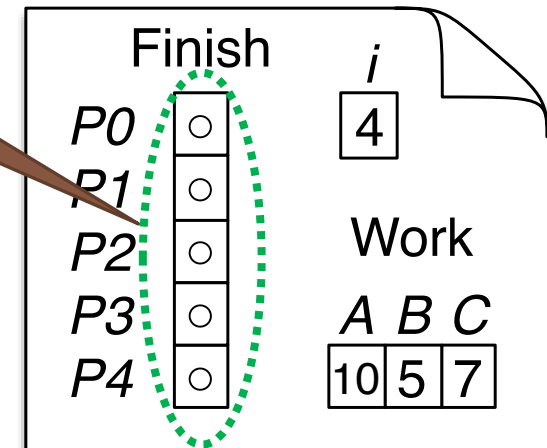
Step

A possible allocation sequence
(P1::P3::P0::P2::P4) exists →
the system is in a **safe state**

- $Work := Work + Allocation[i]$
- $Finish[i] := true$
- System is safe iff.
for all P_j , $Finish[j] == true$

✓ P0
✓ P1
✓ P2
✓ P3
✓ P4

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	3	3	2
P1	2	0	0	1	2	2			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			



Resource Request Algorithm

Request made by process P_i

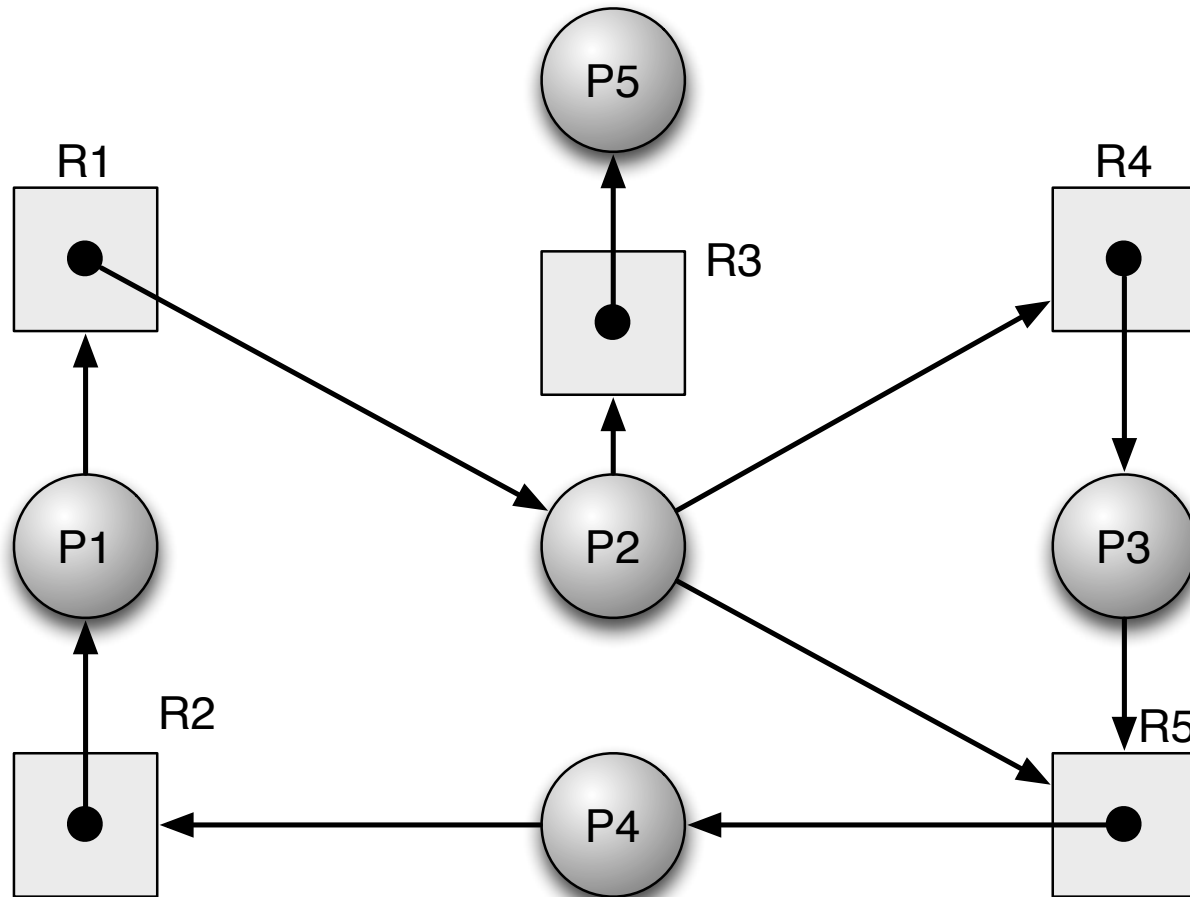
- $Request_Pi[j]$ = # instances requested of type j

Steps

1. If not $Request_Pi \leq Need[i]$ => **reject** (request more than max.)
2. If not $Request_Pi \leq Available$ => P_i must **wait**
3. Try the following new states
 - $Available' := Available - Request_Pi$
 - $Allocation'[i] := Allocation[i] + Request_Pi$
 - $Need'[i] := Need[i] - Request_Pi$
4. Run safety algorithm with State ($Available'$, $Allocation'$, $Need'$)
 - If **safe** => **accept** to give resources
 - If **unsafe** => P_i must **wait**

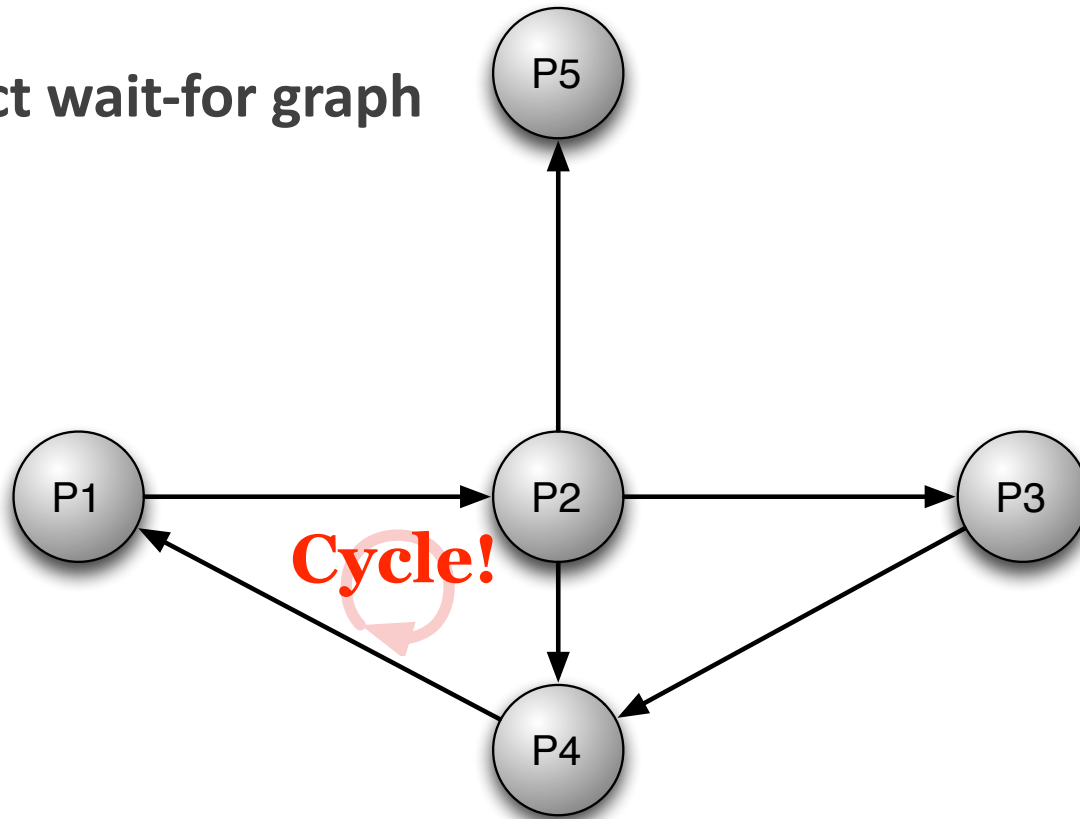
Deadlock Detection

Detection: 1 Inst. / Type



Detection: 1 Inst. / Type (cont.)

Construct wait-for graph



Detection: m Inst. / Type

Variables

- Define $Work := Available$
- Define $Finish[i] := (Allocation[i] == 0)$

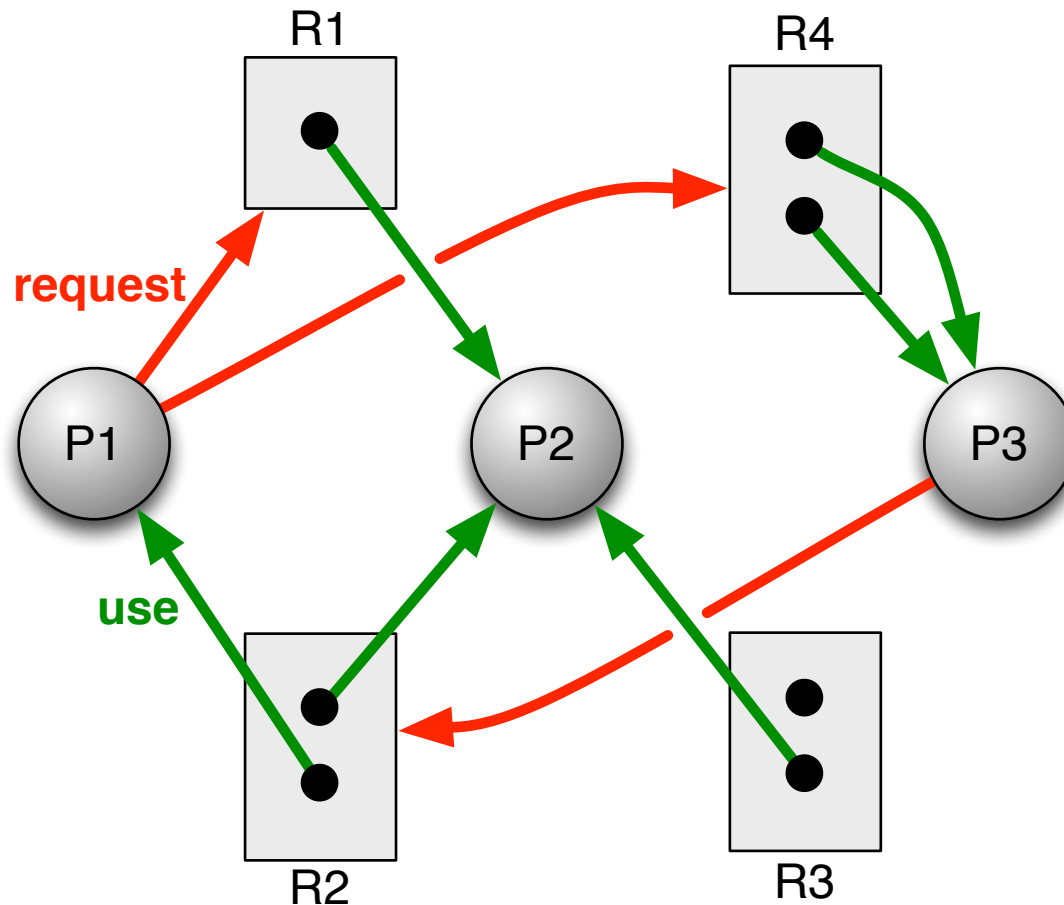
Steps (similar to Banker's safety alg.)

- While exists i s.t. $Finish[i] == false \ \&\& \ Request_P_i \leq Work$
 - $Work := Work + Allocation[i]$
 - $Finish[i] := true$
 - If exists P_j s.t. $(Finish[j] == false) \rightarrow P_j$ is deadlocked

Notes

- Algorithm assumes processes will not ask for more resources
- If they do, then deadlock must be detected again later

Example: m Inst. / Type



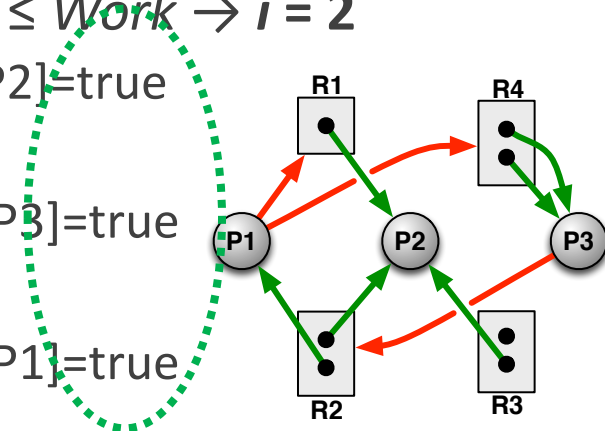
Example: m Inst. / Type (cont.)

Initial conditions

- **Allocation** = [P1: [0 1 0 0] P2: [1 1 1 0] P3: [0 0 0 2]]
- **Request** = [P1: [1 0 0 1] P2: [0 0 0 0] P3: [0 1 0 0]]
- **Work** = [0 0 1 0]
- **Finish** = [false false false]

Steps

- If exists i s.t. $Finish[i] == \text{false} \ \&\& \ Request_P_i \leq Work \rightarrow i = 2$
 - $Work = Work + P2:[1\ 1\ 1\ 0] = [1\ 1\ 2\ 0]$; $Finish[P2] = \text{true}$
- If exists i s.t. ... $\rightarrow i = 3$
 - $Work = Work + P3:[0\ 0\ 0\ 2] = [1\ 1\ 2\ 2]$; $Finish[P3] = \text{true}$
- If exists i s.t. ... $\rightarrow i = 1$
 - $Work = Work + P1:[0\ 1\ 0\ 0] = [1\ 2\ 2\ 2]$; $Finish[P1] = \text{true}$



Recovery From Deadlock

Process termination

- Abort all deadlocked processes
- Abort one process at a time

Resource preemption

- Issues
 - How to select target (process, resource)
 - How to do rollback (e.g., restart)
 - How to avoid starvation

Summary

Deadlock

- Situation when execution cannot continue

Deadlock prevention

- Avoid any of the 4 necessary conditions

Deadlock avoidance

- Safe versus unsafe allocation
- Banker's algorithm: safety alg. + resource request alg.

Deadlock detection

- Similar to Banker's safety algorithm

Next Time

Main memory

Paging

Segmentation